

Introduction to Machine Learning

Presenters



Senior Data Scientist,
Personalization,
Customer Tech



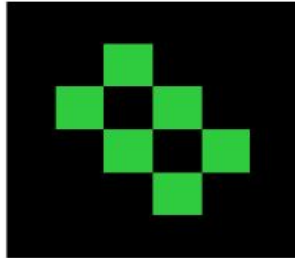
Data Scientist
Data Science Foundations,
Platform

Contents

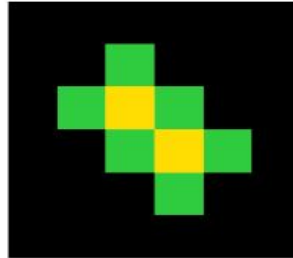
- AI, Machine Learning Overview
- Supervised Machine Learning
 - Regression
 - Classification
- Classification
 - kNN Classifier + code walkthrough
 - Logistic Regression Classifier + code walkthrough
- Data Science Modeling Pipeline
 - Train – Test Splits
 - Feature Pre-processing
 - Diagnosing and Fixing Underfitting/Overfitting

Humans are great at

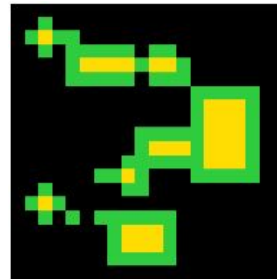
Abstraction
and
Reasoning



Input



Output



Humans are great at ...

Learning tasks.

Example: Humans learning to drive a car **vs** training a Self Driving Car.



Humans are good at

Planning:

- A birthday party
- A vacation
- Business strategies
- Navigating Traffic

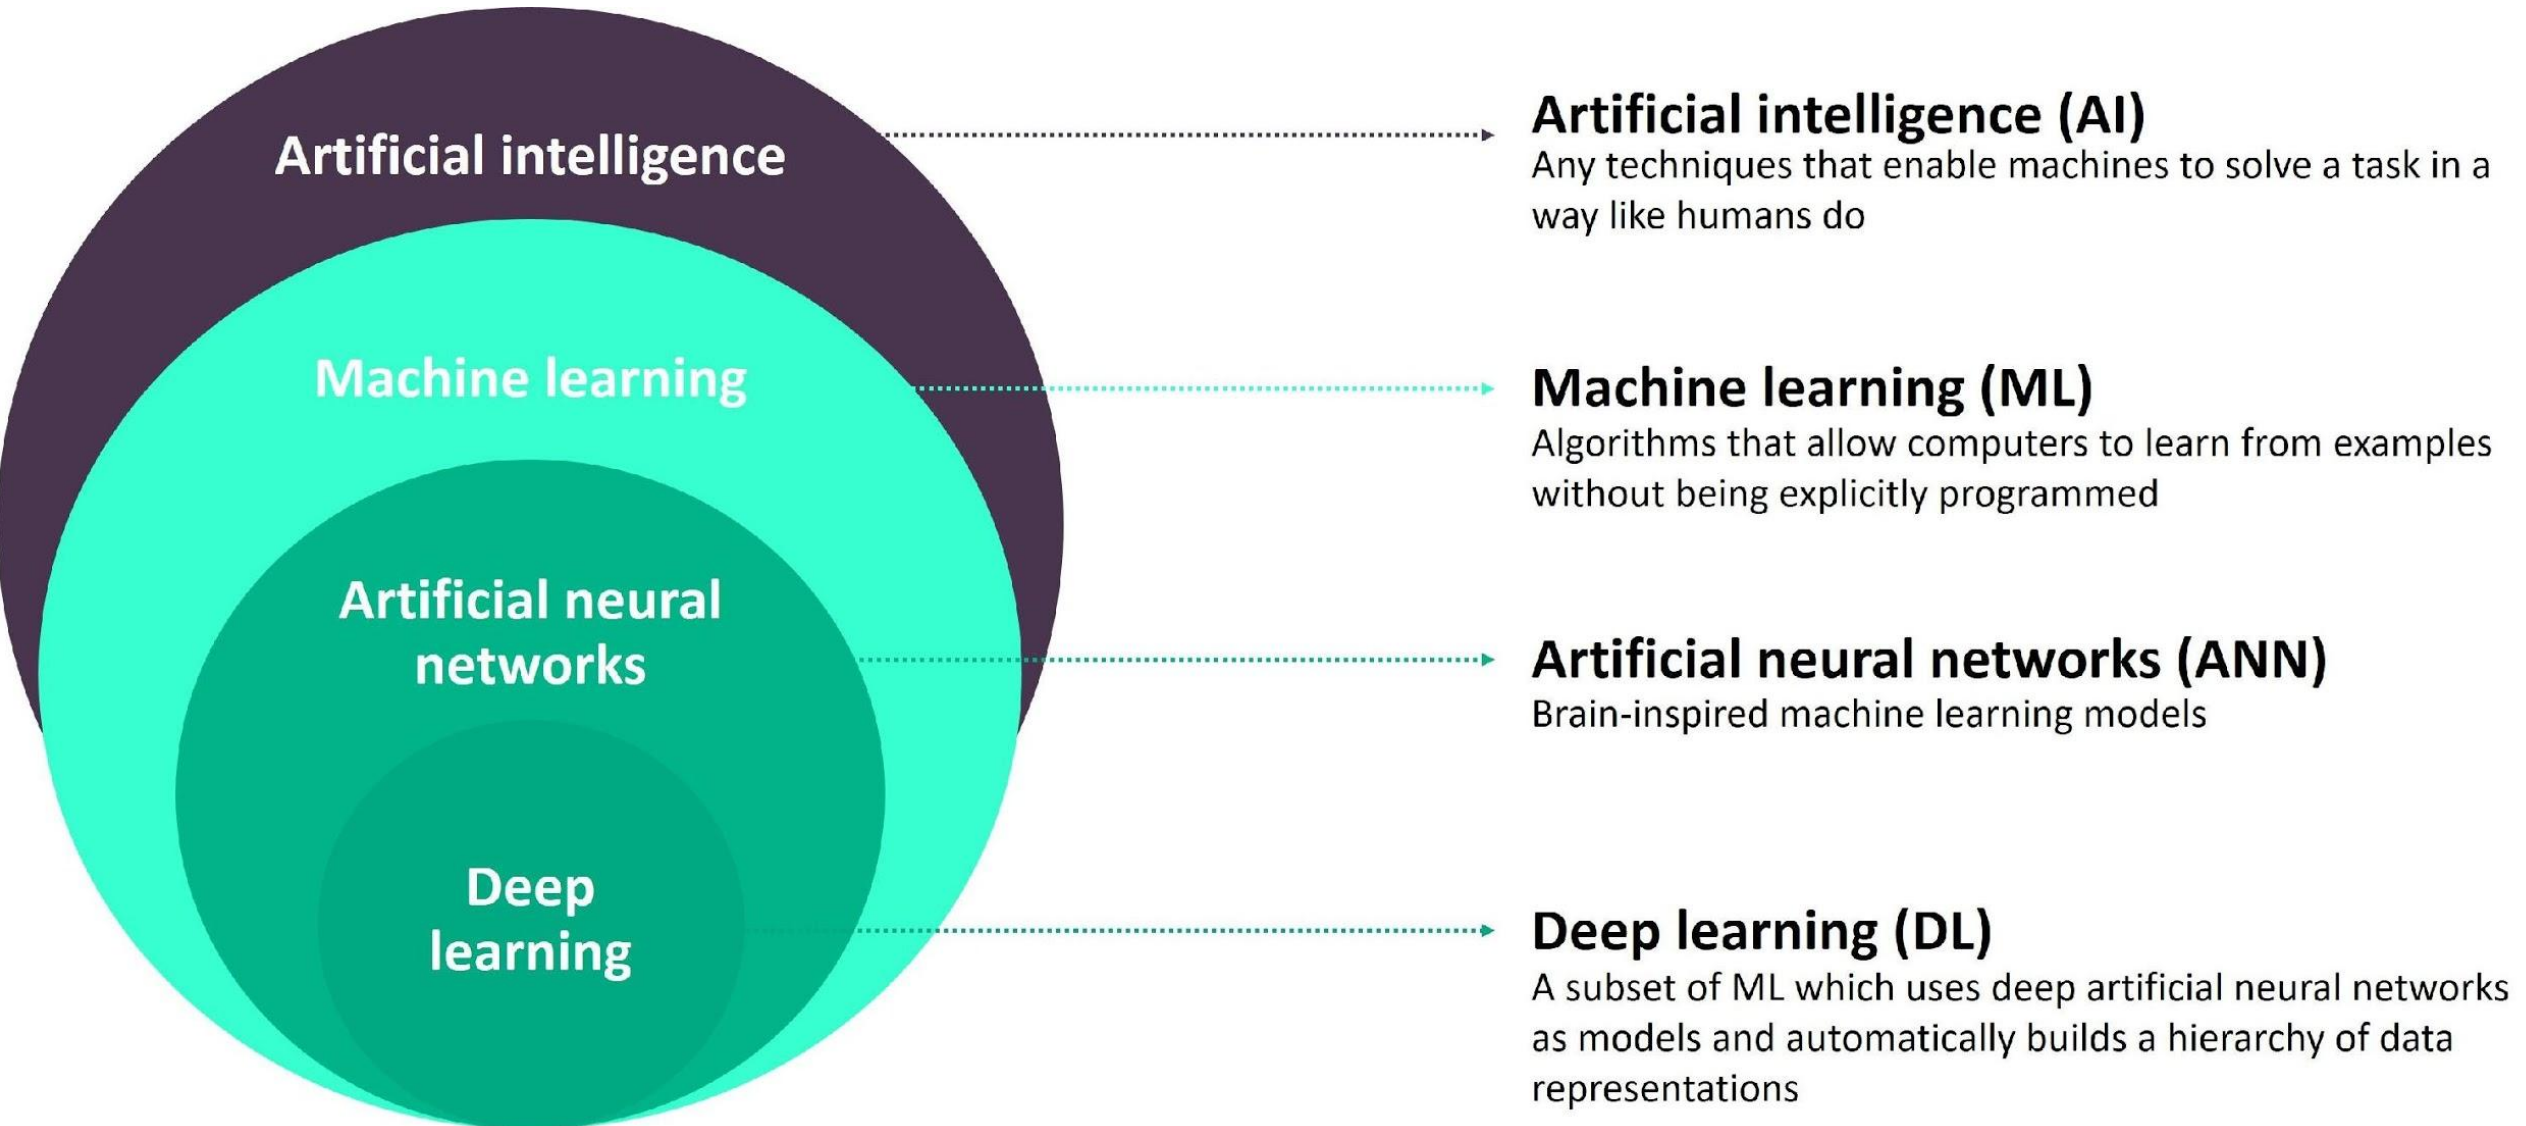
Natural Language Understanding

- Complex sentences (Legal)
- Short sentences (Twitter)
- Sarcasm
- Humor
- Emotion
- Innate understanding of the world

Visual Perception

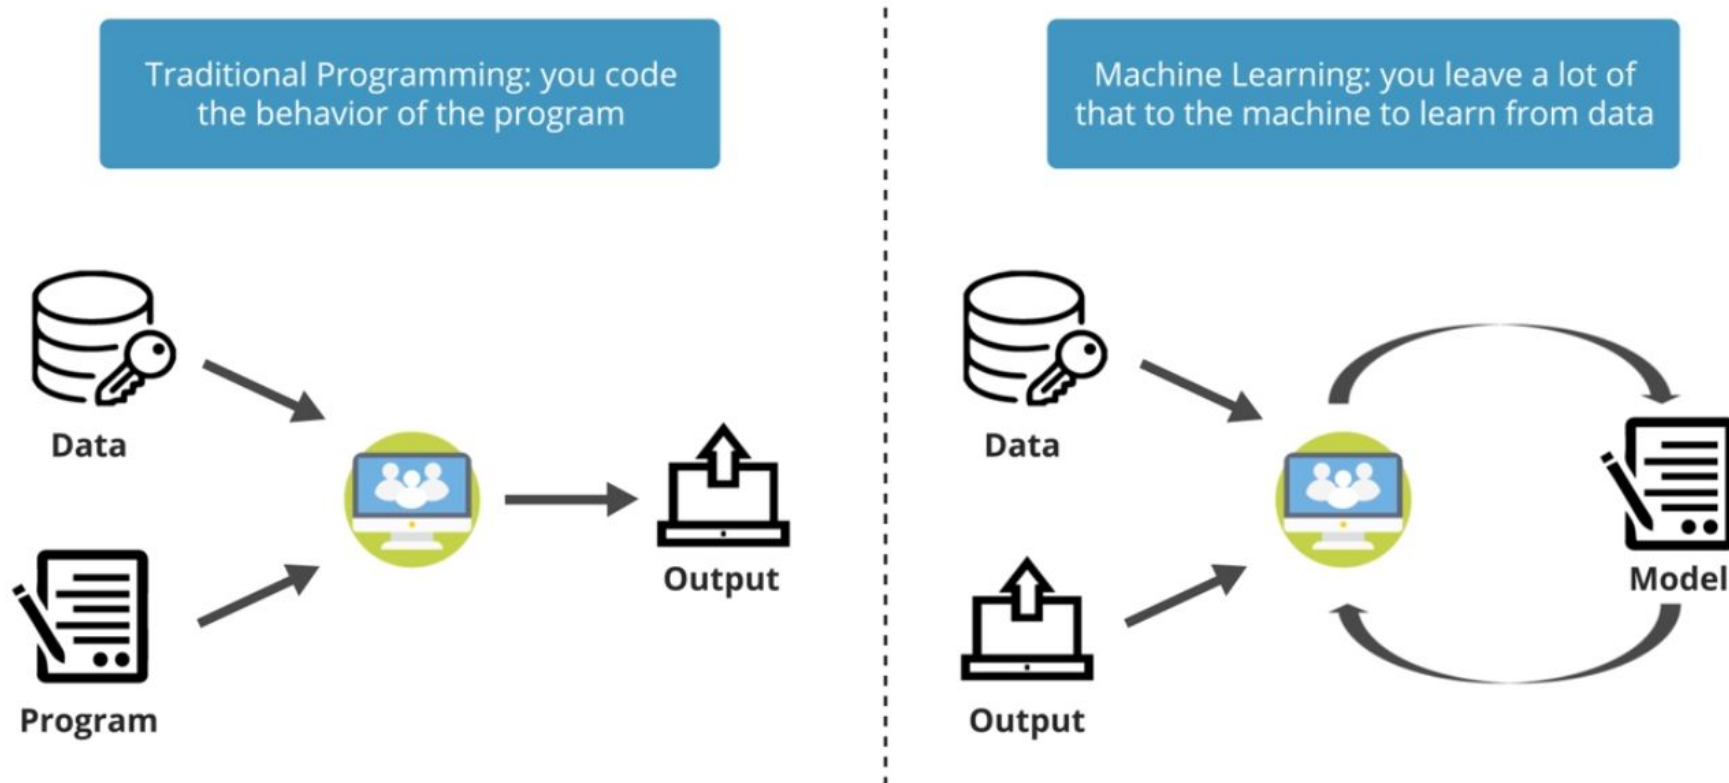
- Scene Understanding
- Depth Perception
- Object Recognition

Artificial Intelligence



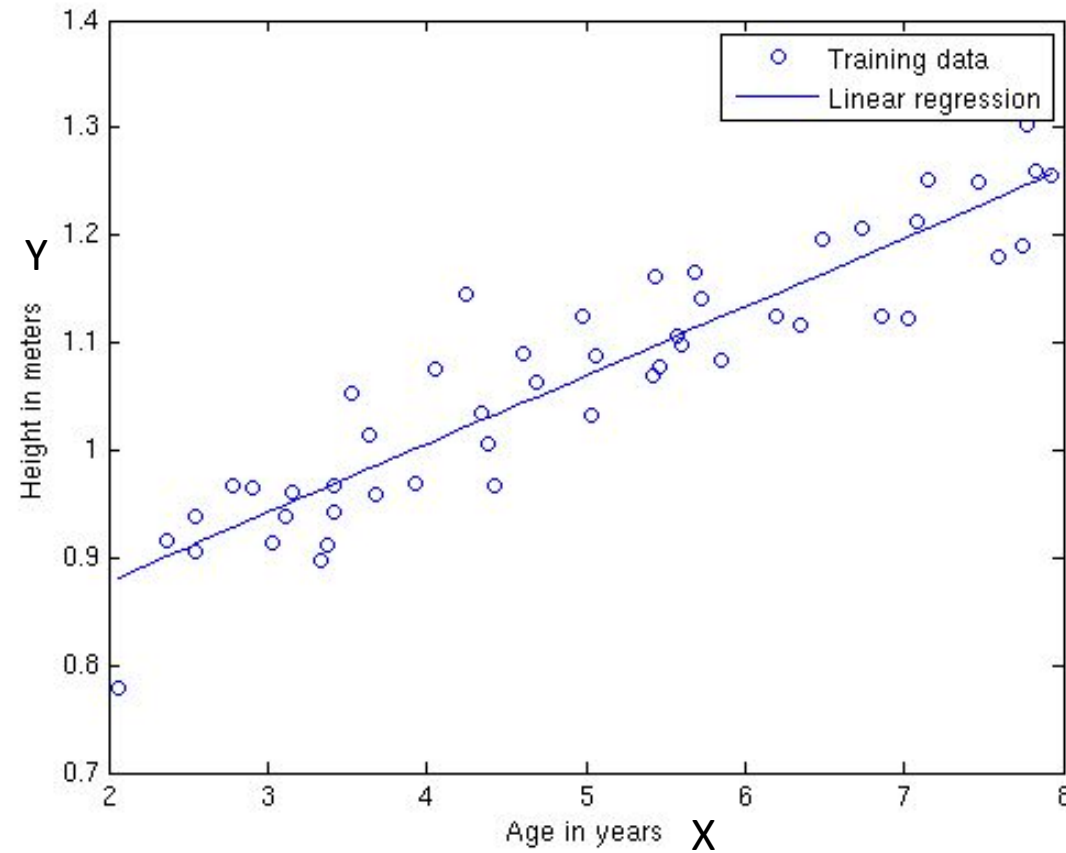
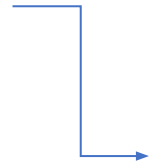
Supervised Machine Learning

Traditional Approach vs. Machine Learning Approach

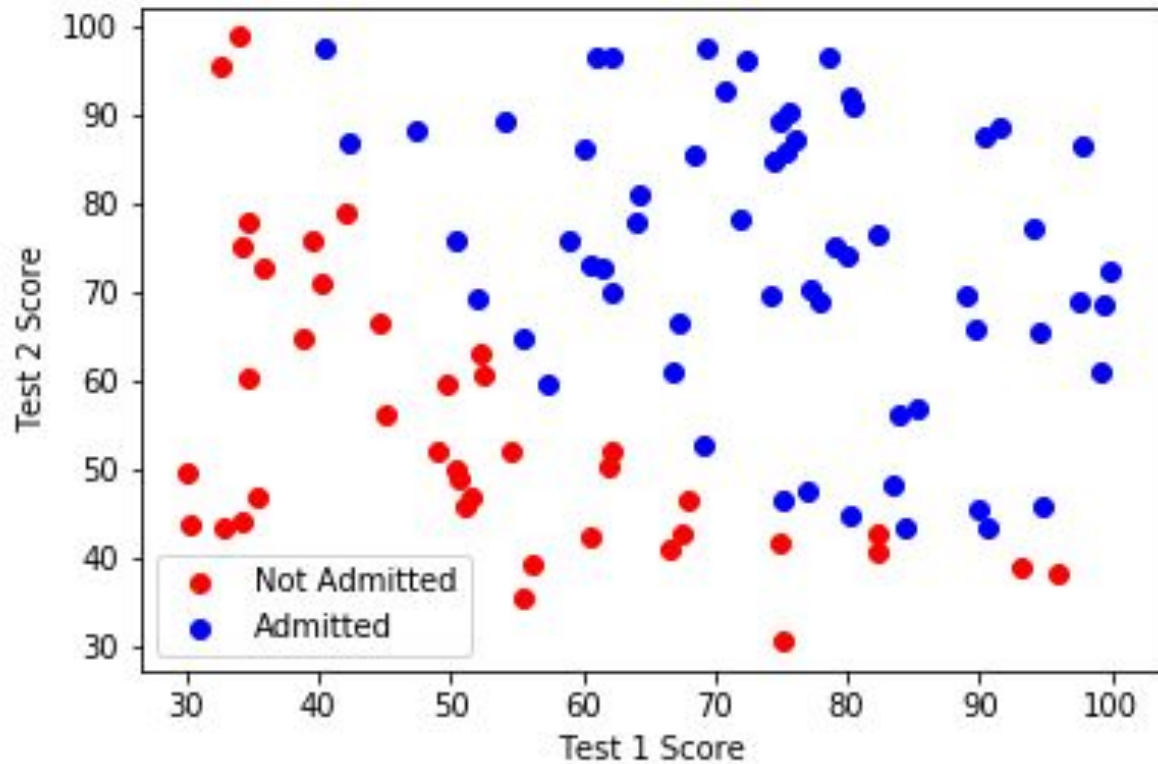


Supervised Learning – Regression Problem

**Continuous
Valued Target**

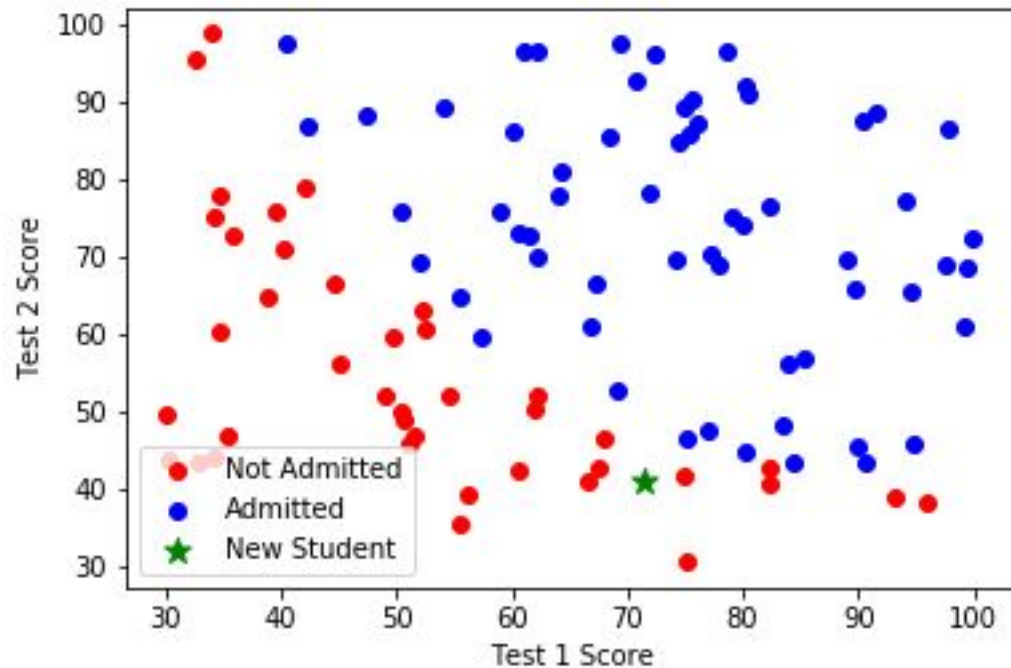


Predicting a Student's Admission



Historical Data OR
Training Data

Predicting a Student's Admission

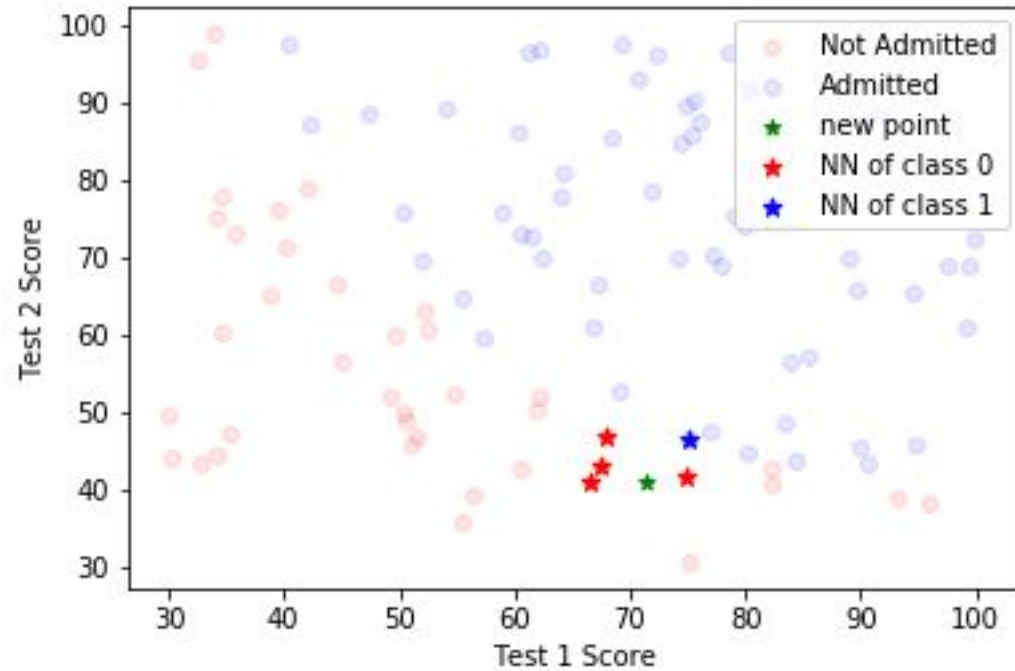


**Binary
Classification Task:**

Predict whether a
new student will
secure admission.

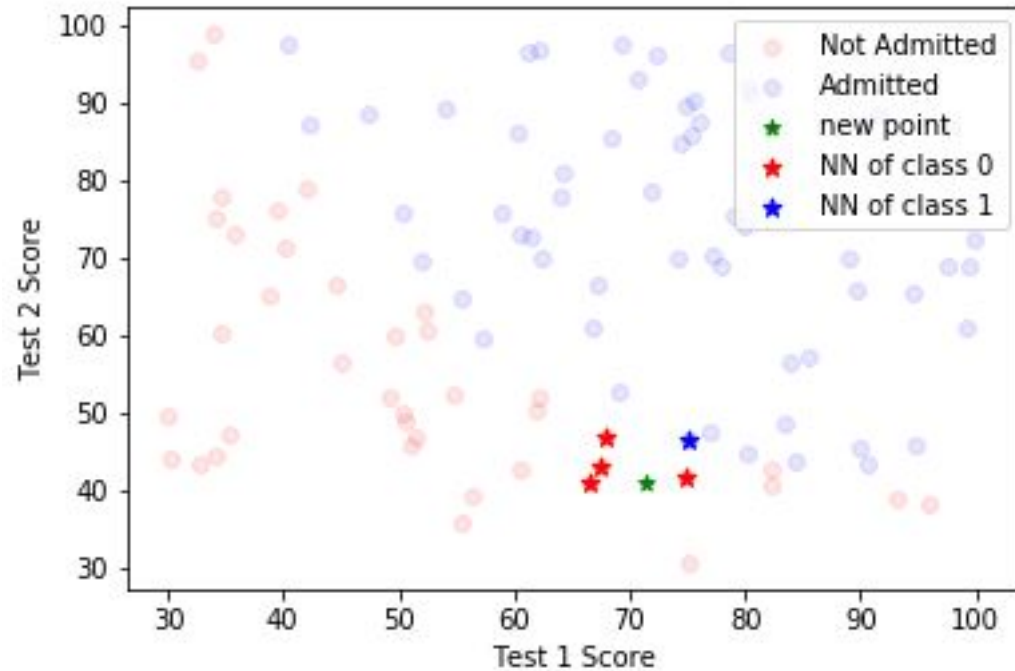
New student will not
be present in
historical data

k-Nearest Neighbours (kNN)

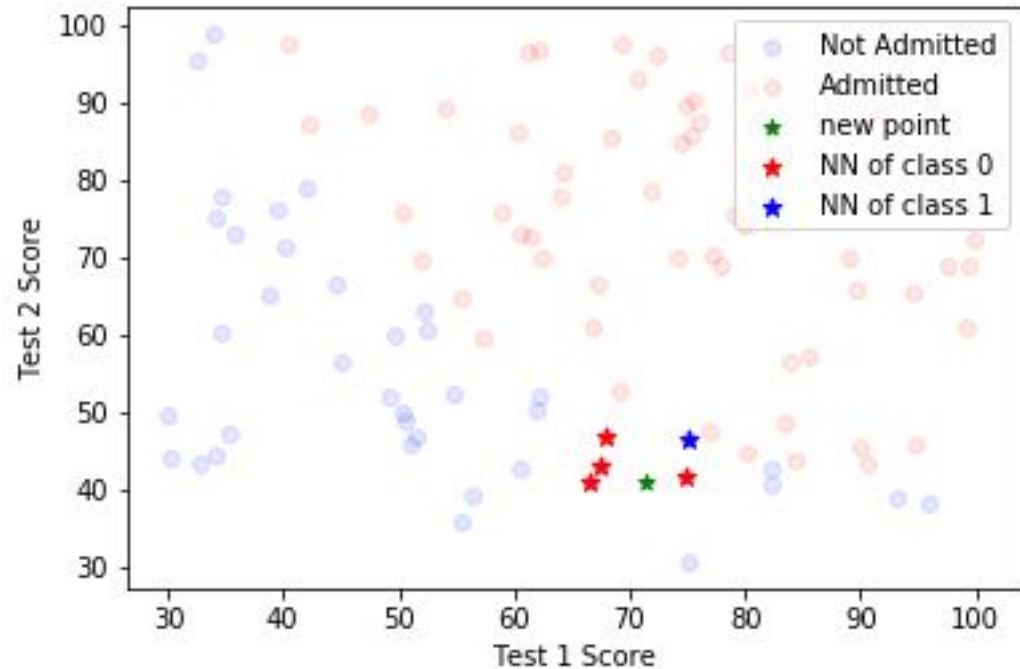


Indicate the
closest points
in historical
data to the
new point

kNN Classification – Majority Voting



kNN Classification – Probability Output



4 ★
1 ★
K = 5 ★

Let $\hat{y}=1$ be the event
“Student is Admitted”

$$Prob(\hat{y} = 1 | x) = \frac{n_1}{K} = \frac{1}{5} = 0.2$$

Where n_1 is the
number of nearest
neighbors of class 1

kNN – The Math

$$x = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}$$

x is a new
observation
column vector

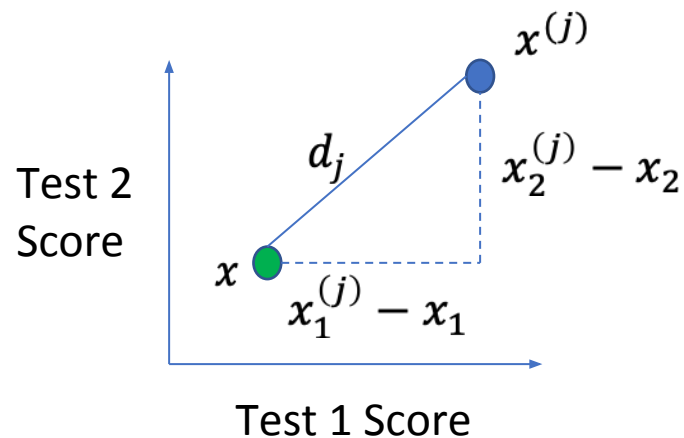
$$X = \begin{bmatrix} x^{(1)T} \\ x^{(2)T} \\ x^{(M)T} \end{bmatrix} = \begin{bmatrix} x_1^{(1)} & x_2^{(1)} \\ x_1^{(2)} & x_2^{(2)} \\ x_1^{(M)} & x_2^{(M)} \end{bmatrix} \quad y = \begin{bmatrix} y^{(1)} \\ y^{(2)} \\ y^{(M)} \end{bmatrix}$$

Design Matrix

Where,

x_1 = Test 1 Score

x_2 = Test 2 Score



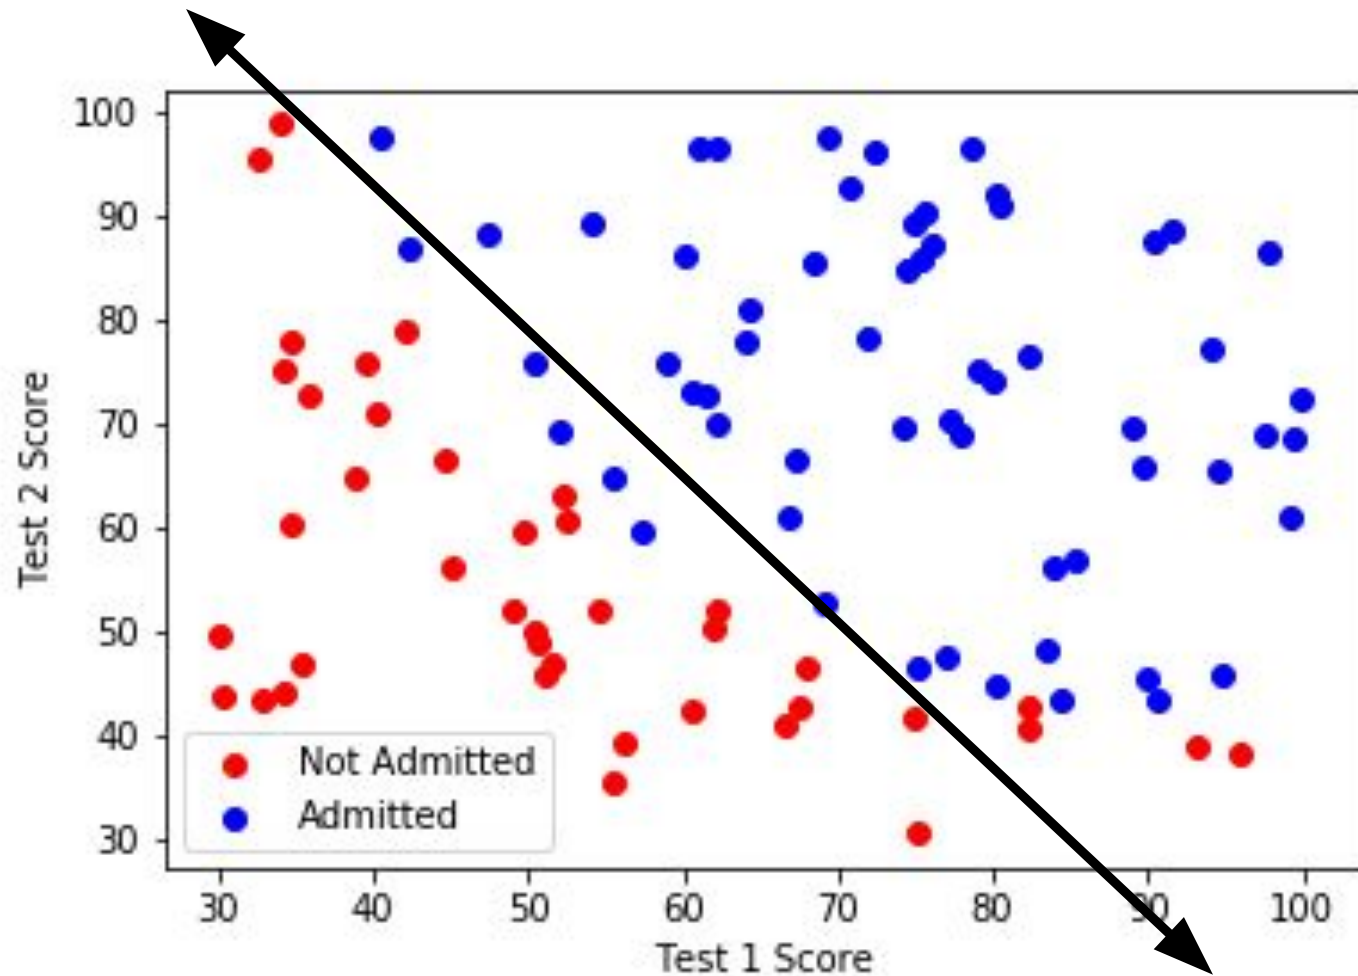
$$d_j = \sqrt{\left(x_1^{(j)} - x_1\right)^2 + \left(x_2^{(j)} - x_2\right)^2}$$

**From Pythagoras'
Theorem**

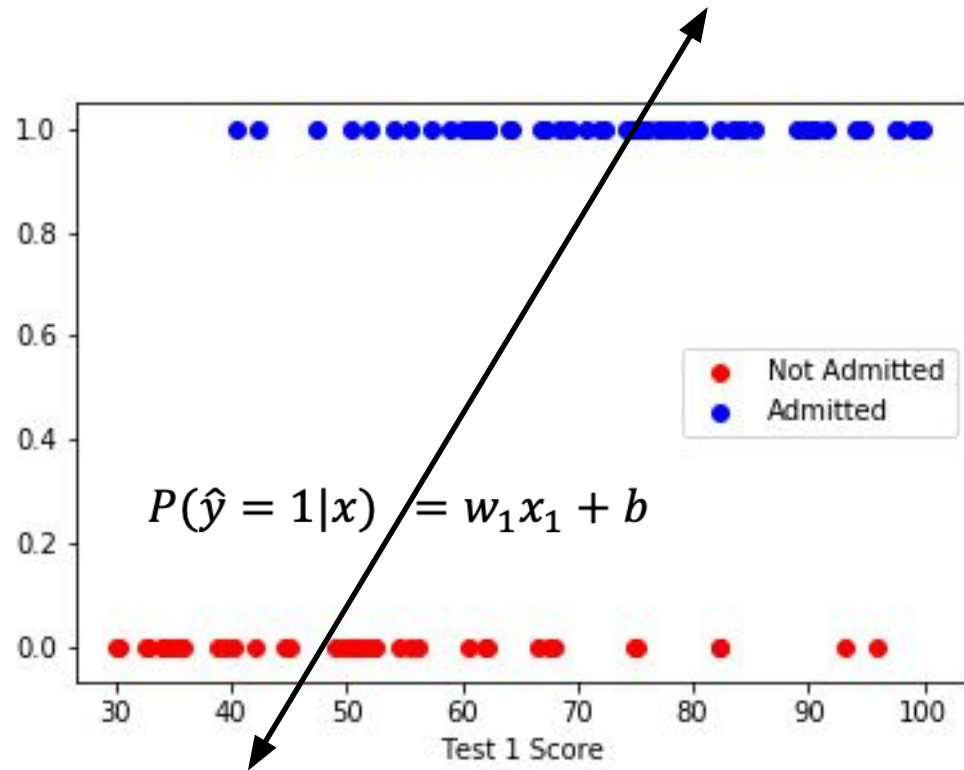
$$NN(x, X) = \arg \min_{j \in \{1, M\}} d(x, x^{(j)})$$

QnA & Code Walkthrough - kNN

Logistic Regression



First, a linear probability model



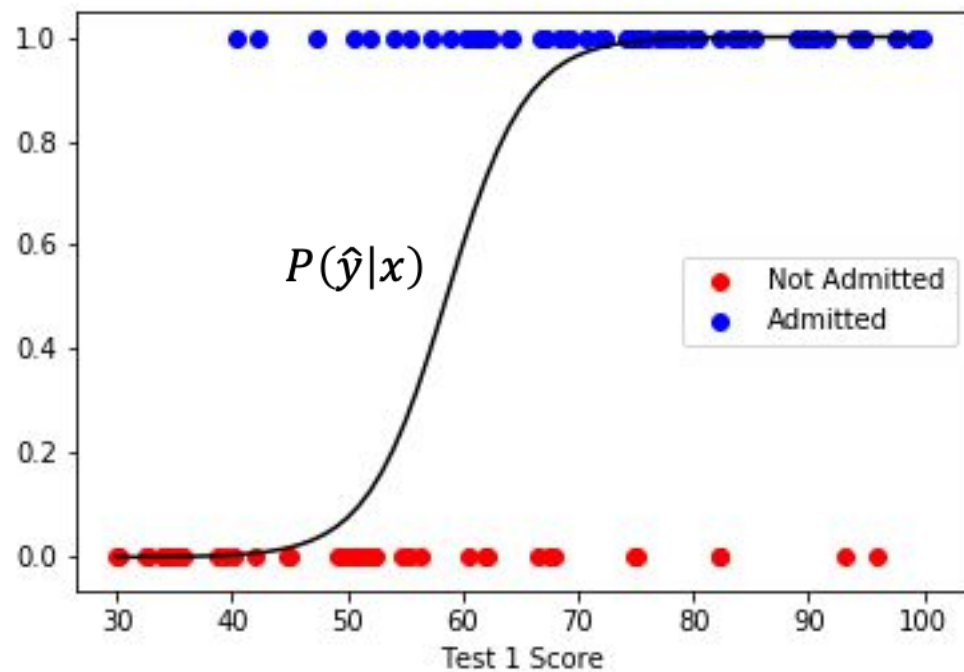
Estimate the coefficients w_1 and b through Least Squares Regression

$$w = (X^T X + \lambda I)^{-1} X^T Y$$

A Major Problem with this Model:

Non-sensical probabilities ($P > 1$ or $P < 0$)

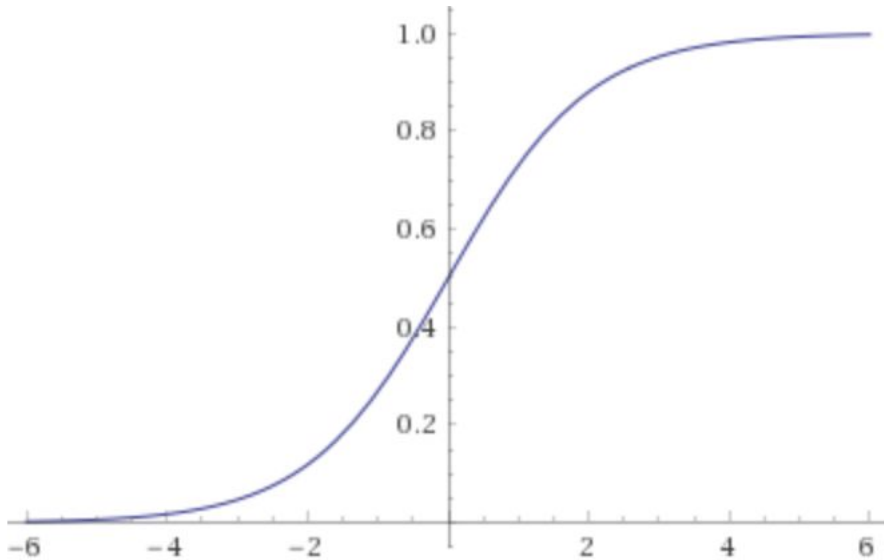
Would be nice, if ...



Output bounded
between 0 and 1.

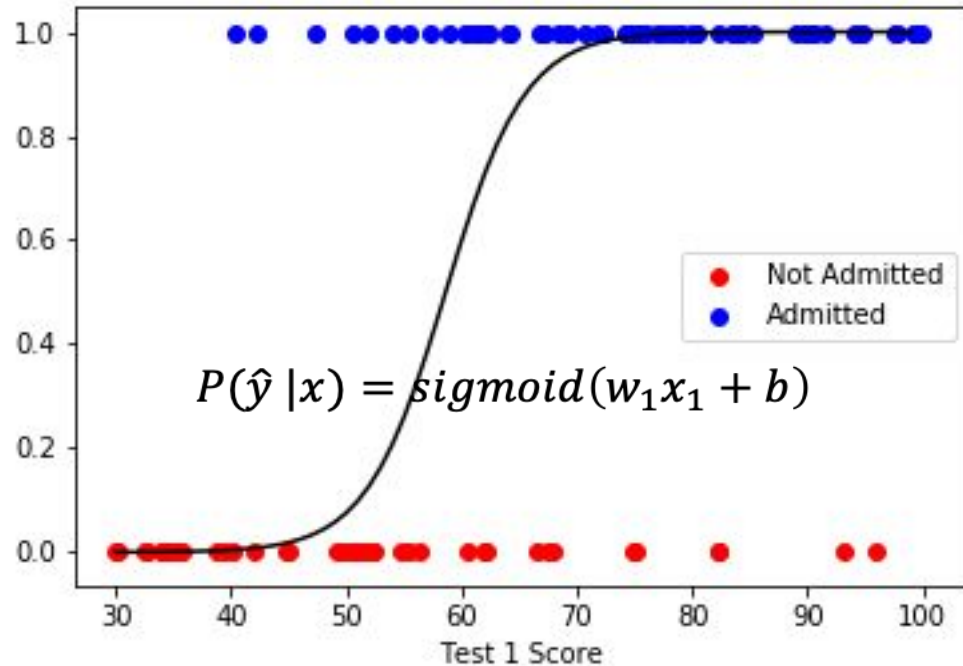
Sigmoid/Logistic Function

$$\text{sigmoid}(x) = \frac{1}{1 + e^{-x}}$$



- The sigmoid function squishes numbers to (0,1)

Logistic Regression



[Notebook on Github](#)

[Scikit-learn implementation of
Logistic Regression](#)

Unknowns

$$P(\hat{y} = 1 | x) = \text{sigmoid}(\mathbf{w}_1 x_1 + b) = \sigma(w_1 x_1 + b) \\ = \frac{1}{1 + e^{-(w_1 x_1 + b)}}$$

General Form:

$$P(\hat{y} = 1 | x) \\ = \sigma(\mathbf{w}_1 x_1 + \mathbf{w}_2 x_2 + \dots + \mathbf{w}_N x_N + b) \\ = \sigma(\mathbf{w}^T \mathbf{x} + b)$$

$$\begin{matrix} \boxed{w_0} & \boxed{w_1} & \boxed{w_2} & \dots & \boxed{w_N} \\ \mathbf{w}^T \end{matrix} \quad \begin{matrix} \boxed{x_0} \\ \boxed{x_1} \\ \boxed{x_2} \\ \dots \\ \boxed{x_N} \end{matrix} \quad + \quad \boxed{b}$$


Q & A

Loss Minimization


- Most of machine learning involves some form of loss minimization
- Loss indicates how bad our predictions are.
- Let \hat{y} be our predicted probability of admission of a student x , and y be the true label.
- Then, $L = f(\hat{y}, y)$ is the loss we incur, where $f(., .)$ is called the loss function

Some Examples of Loss Functions

- **0-1 Loss:** $\hat{y} = 0, y = 1 \Rightarrow L = 1$
- **Squared Loss:** $\hat{y} = 0.2, y = 1 \Rightarrow L = (y - \hat{y})^2 = 0.8^2 = 0.64$
- **Log Loss / Binary Cross Entropy Loss:** This is the loss function used in Logistic Regression.
- **Important:** Loss is a function of the weights and not the data.
 - $$\begin{aligned} L &= f(\hat{y}, y) \\ &= f(h(x; \mathbf{w}, \mathbf{b}), y) \end{aligned}$$


$$h(x; \mathbf{w}, \mathbf{b}) = \sigma(\mathbf{w}^T \mathbf{x} + b)$$

Average Loss Function

$$L = \frac{1}{M} \sum f(\hat{y}_i, y_i)$$

$$= \frac{1}{M} \sum f(h(x_i; w, b), y_i)$$

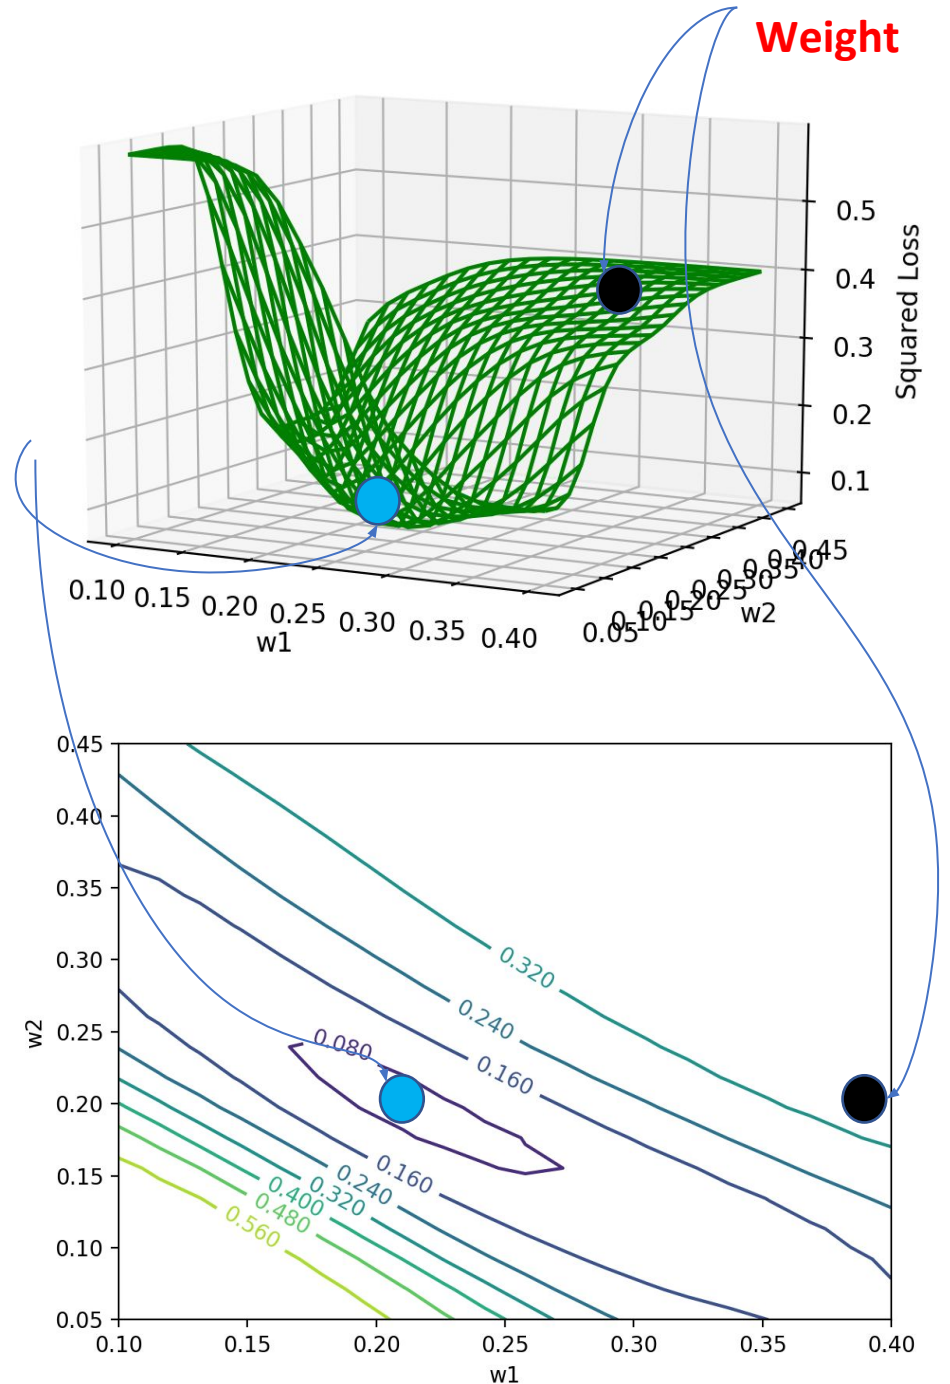
- $\hat{y}_i = h(x_i; w, b)$ is the prediction for the i th sample in dataset
- L is the average loss.

Plotting Loss Functions

- Plot average loss for every possible $w = [w_1, w_2]$
- Plotting this will give us a 3D plot. (top figure)
- Plotting equi-loss surface will give us a contour plot (bottom figure)

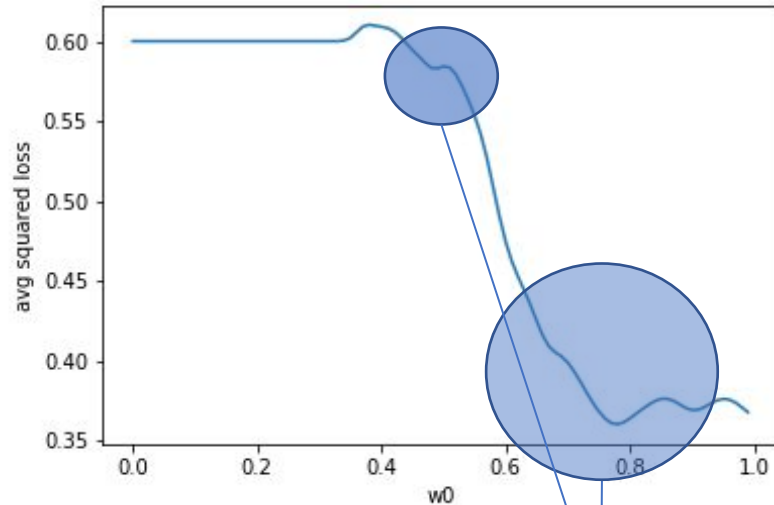
Optimal
Weights

Random
Weight



Appropriate Loss Function for Logistic Regression

Squared Loss

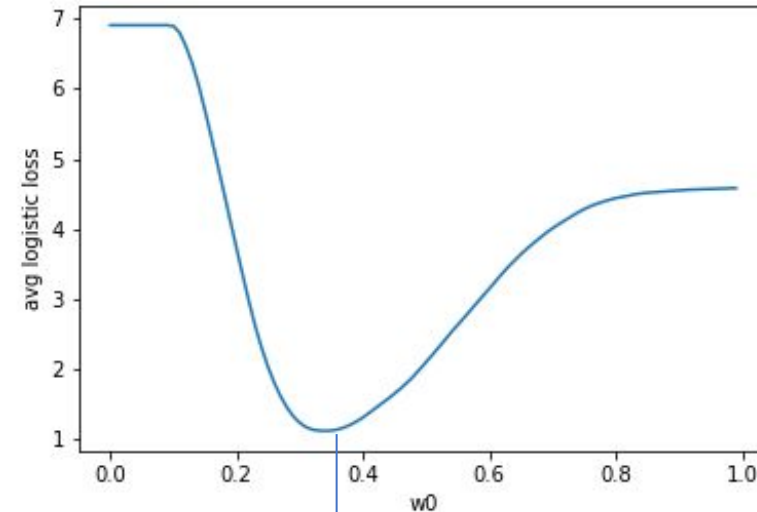


Word of Caution:
Squared Loss
isn't bumpy
when used in
Linear
Regression.

Many Local Minima !

Fixed $w_1 = -0.41$
And $b = -21$

Log Loss



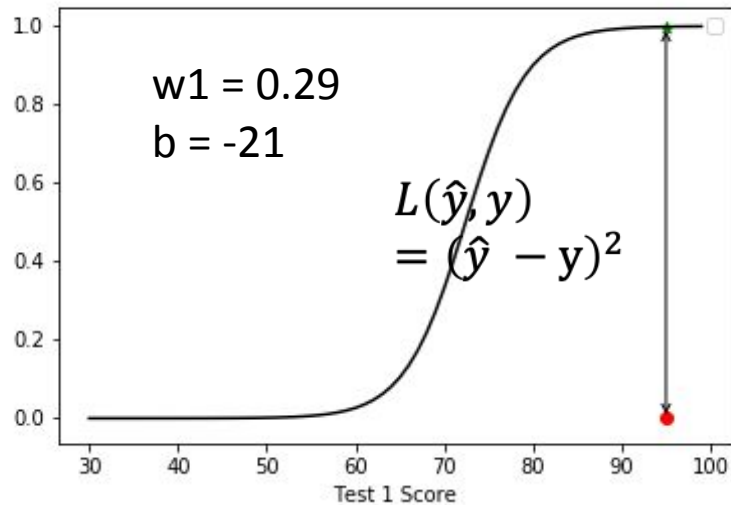
1 Global Minima

A Loss Function for Logistic Regression

**Squared
Loss**

Observe:

The max loss can be only 1

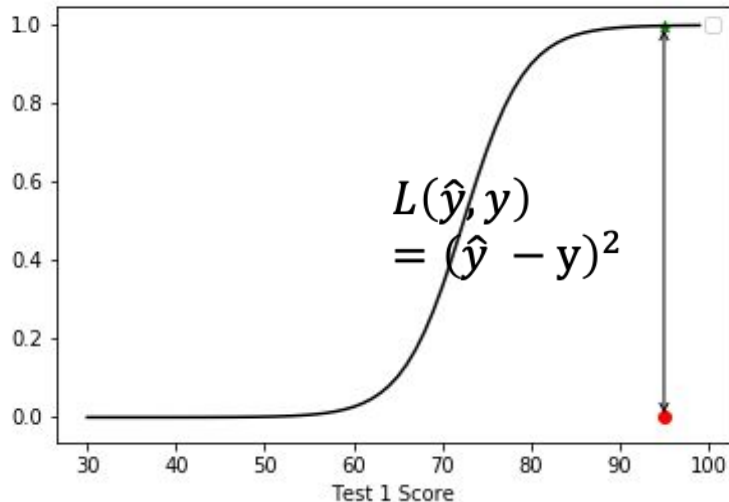


A Loss Function for Logistic Regression

Squared Loss

Observe:

The max loss can be only 1



What if ?

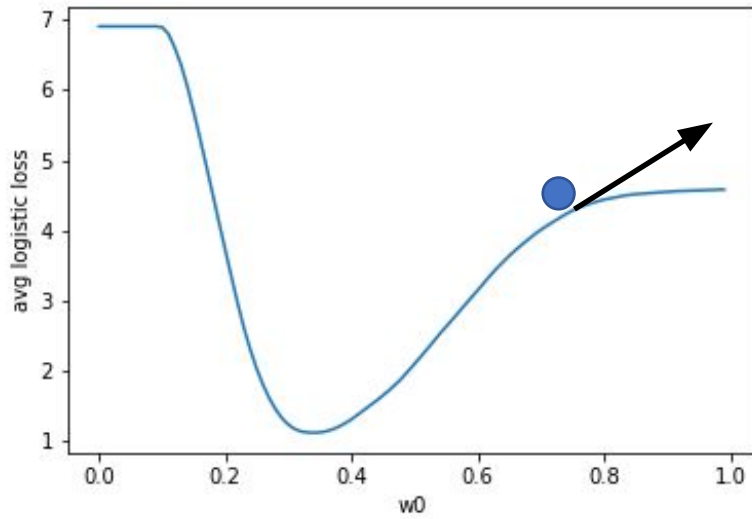
- When $y = 1$ and $\hat{y} = 0$, $f = \infty$
- When $y = 0$ and $\hat{y} = 1$, $f = \infty$

$$f(y, \hat{y}) = \begin{cases} -\ln \hat{y}, & \text{if } y = 1 \\ -\ln(1 - \hat{y}), & \text{if } y = 0 \end{cases}$$

$$f(y, \hat{y}) = -y \ln \hat{y} - (1 - y) \ln(1 - \hat{y})$$

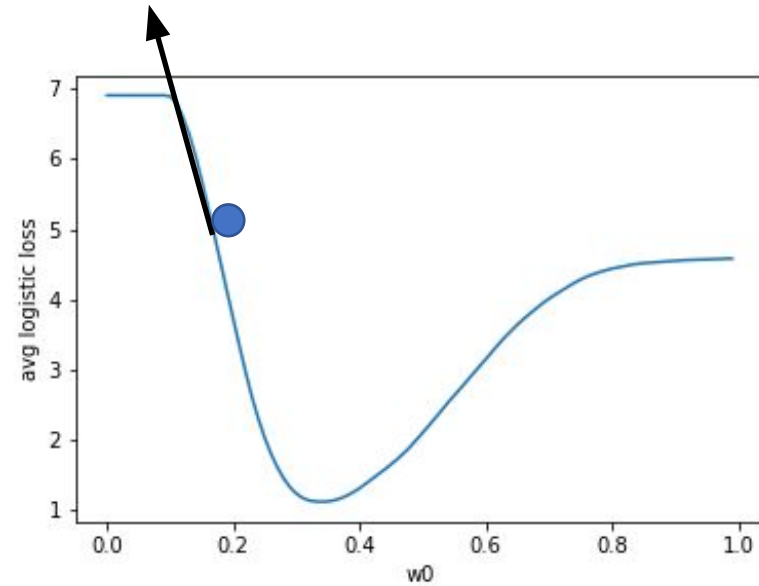
Cross Entropy Loss

Gradient

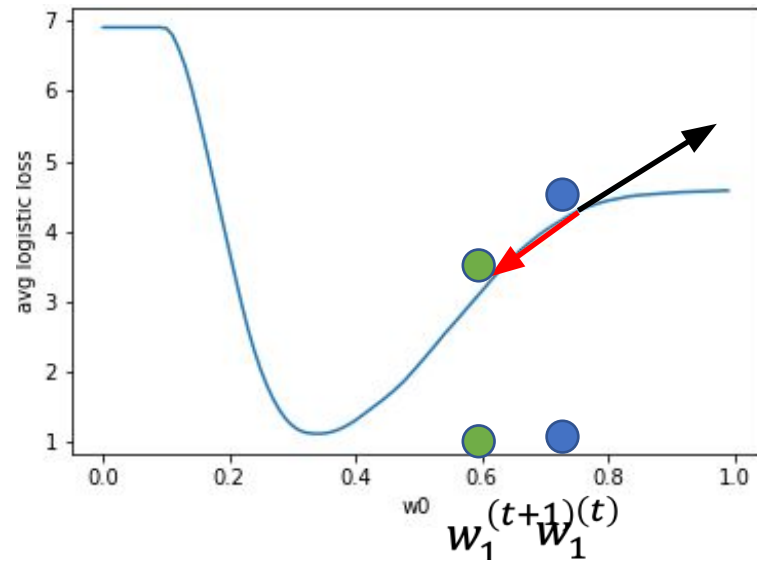


$$Tangent = \frac{\partial L}{\partial w_1}$$

Black arrow
magnitude and
direction of tangent
(gradient)



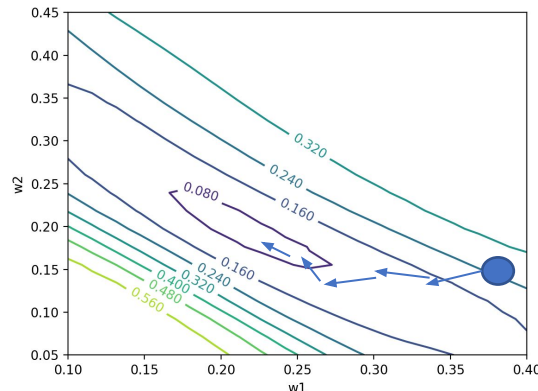
Update Mechanism - Gradient Descent (Roll down the hill)



- Compute the gradient (tangent) at the current parameter values
- Step in the opposite direction

$$w_i^{(t+1)} := w_i^{(t)} - \eta_t \frac{\partial L}{\partial w_i^{(t)}}$$

Updated coeff Learning Rate Gradient



$$\begin{aligned} w_i^{(t+1)} &= w_i^{(t)} - \eta \frac{\partial L}{\partial w_i^{(t)}} \\ &= w_i^{(t)} - \eta \frac{1}{M} \sum (\hat{y}_i - y_i) x_i \end{aligned}$$

Putting it all Together

1. **Given:** Dataset $D = \{(x_1, y_1), \dots, (x_M, y_M)\}$
2. **Initialize:** coefficients w of model randomly
3. $L(w) = \frac{1}{M} \sum_i f(y_i, \hat{y}_i)$
4. For all coefficients w_j :
 1. $g_j = \frac{1}{M} \frac{\sum \partial f(y_i, \hat{y}_i)}{\partial w_j} = \frac{1}{M} \sum (y_i - \hat{y}_i) x_j$
5. For all coefficients:
 1. $w_j = w_j - \eta g_j$
6. Repeat 3-5 till change in loss is negligible

1 Pass /
Epoch

Updates have to be
simultaneous

Computing gradients
over the full dataset
might be expensive.

- Compute over mini batches of data instead (mini-batch gradient descent)

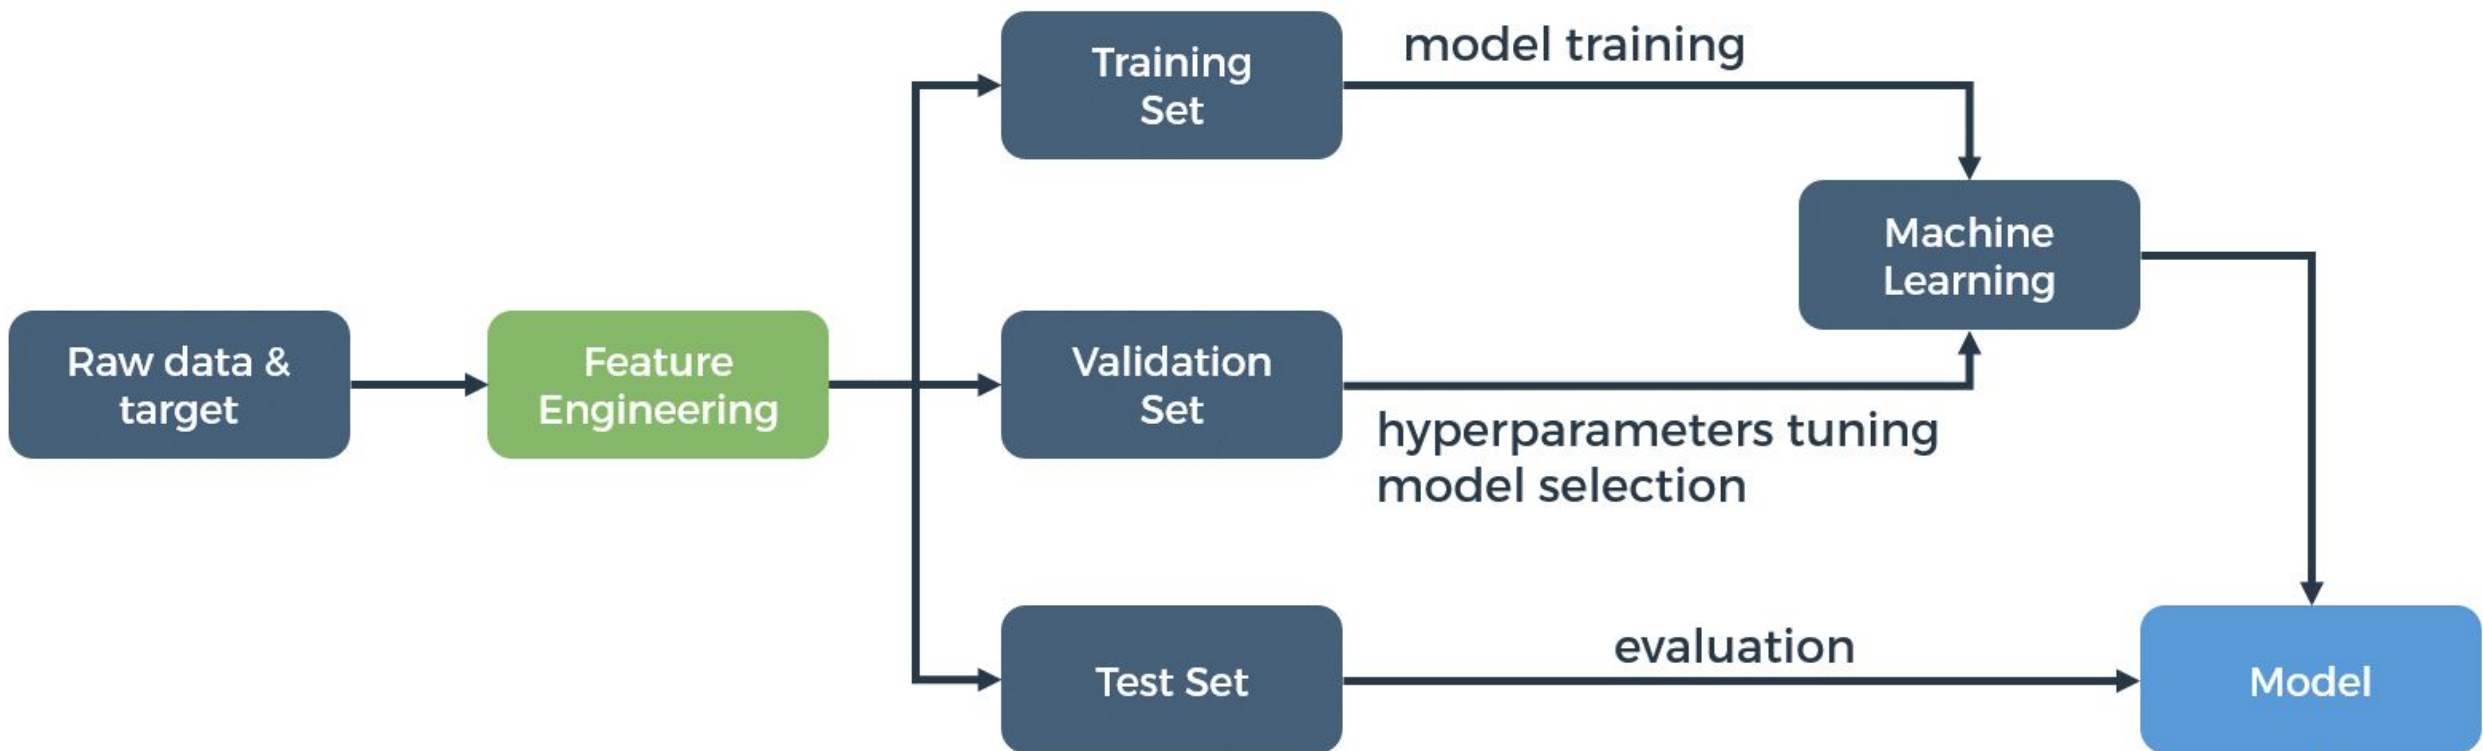
Scikit-learn has an efficient
Logistic Regression
implementation:

Choose a learning rate just
high enough so that training
doesn't diverge (i.e. losses
don't increase with steps)

Q & A

Data Science Pipeline

TRAINING

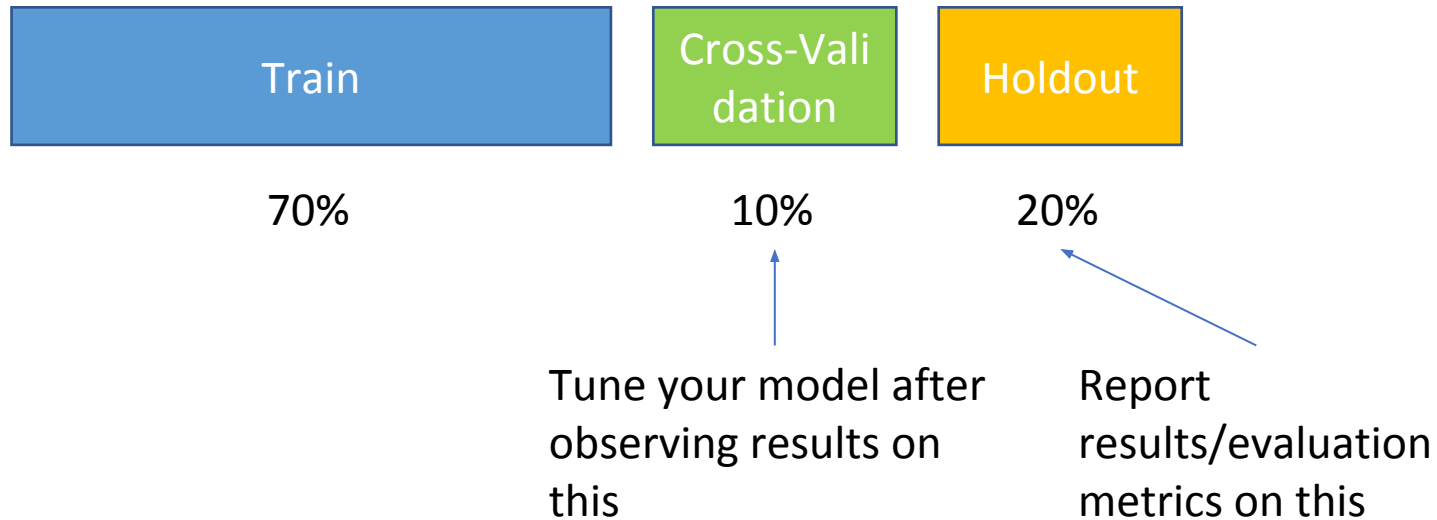


PREDICTING



Train – Test Splits

For Large Data:



For small data, look at K-Fold splits.

Split Strategies:

- 1) Random (70-10-20) split
- 2) Out of time cross validation and holdout

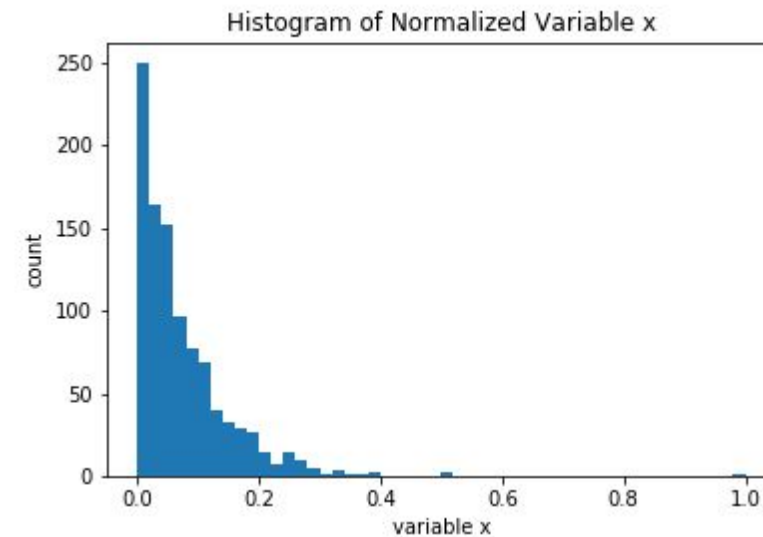
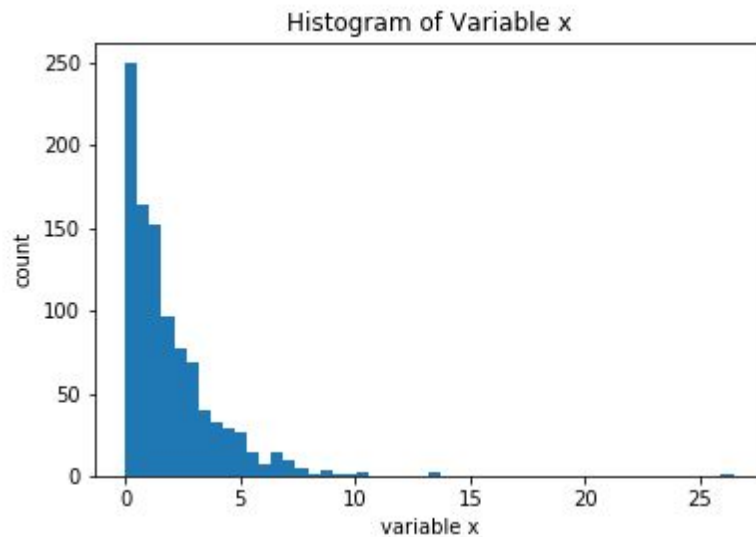
Scikit-learn has implementations of

- Train-Test Split
- K-Fold

Feature Pre-Processing for Continuous variables

- Scale your features to small values around 0
 - Min-Max Scaler $x := \frac{x - x_{mn}}{x_{mx} - x_{mn}}$ [[Scikit Learn MinMaxScaler](#)]

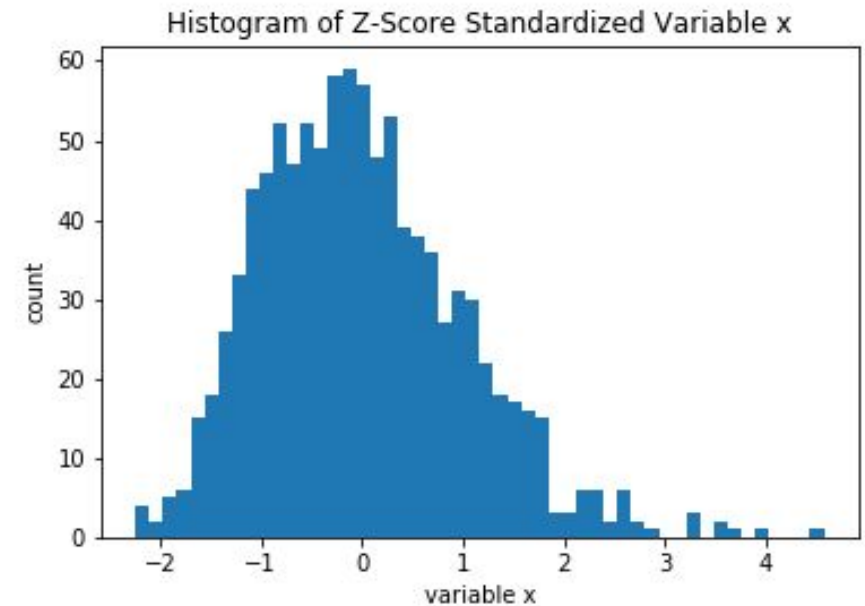
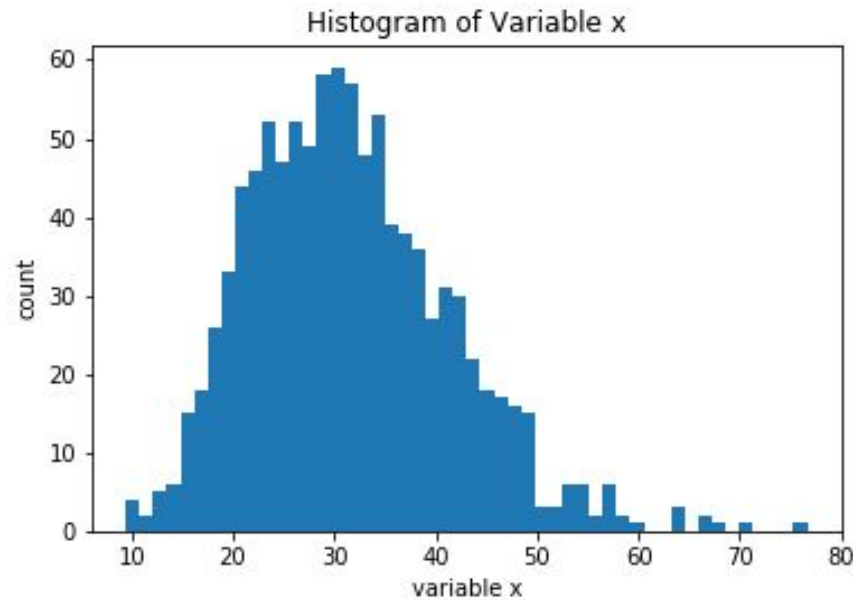
Scaling your features will help the gradient descent converge faster



Feature Pre-Processing for Continuous variables

- Scale your features to small values around 0
 - Z-Score $x := \frac{x - \mu}{\sigma}$ [[Scikit Learn Standard Scaler](#)]

Scaling your features will help the gradient descent converge faster



Feature Pre-Processing for Discrete/Categorical Variables

A categorical variable can take K discrete values with no notion of ordering or rank between them.

Device Type:

Mobile
Laptop
Tablet

1-hot
encode



1	0	0
0	1	0
0	0	1
Mobile	Laptop	Tablet

Scikit-learn

- One Hot Encoder

For a variable with large K, look at other techniques:

- Hashing Trick
- Target Statistics

Model Evaluation Measure - Accuracy

S.No	Predicted Label	Ground Truth Label
1	1	1
2	0	0
3	0	1
4	1	0
5	0	0
6	0	0

$$Accuracy = \frac{N_{correct}}{N} = \frac{4}{6} = 66.67\%$$

Imagine a scenario where 99% of the ground truth labels are 0s

A classifier which labels every example as a 0, will also have 99% accuracy !

Using Accuracy for imbalanced classes will be misleading !!

Model Evaluation Measures

	Actual Positive	Actual Negative
Predicted Positive	tp (True Positives)	fp (False Positives)
Predicted Negative	fn (False Negatives)	tn (True Negatives)

Confusion Matrix

$$\text{Accuracy} = \frac{tp + tn}{tp + fp + fn + tn}$$

$$\text{Recall} = \frac{tp}{tp + fn}$$

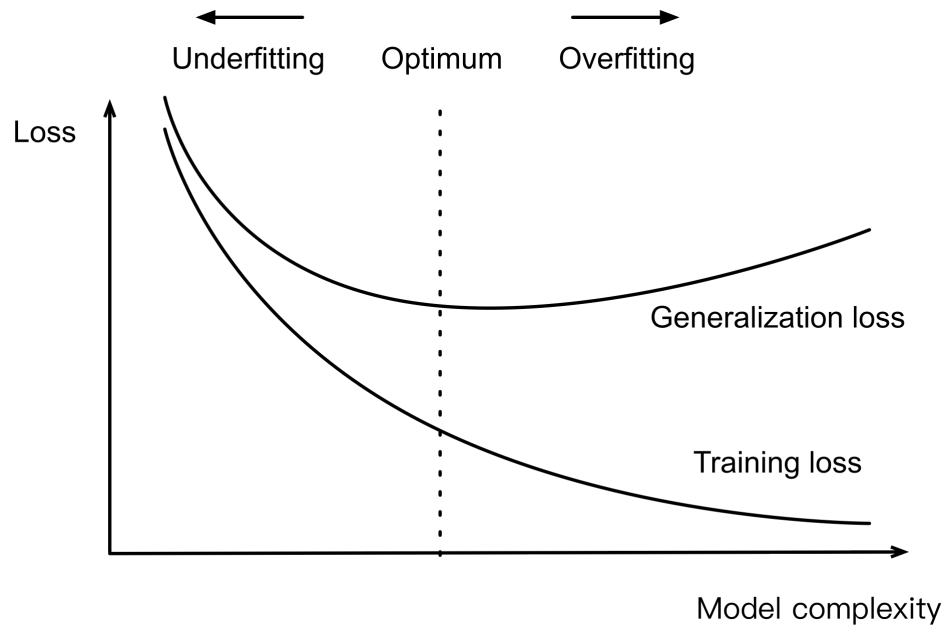
$$\text{Precision} = \frac{tp}{tp + fp}$$

In the case of Ad-Click Prediction:

- If optimizing for reach, then tune for recall
- If optimizing for ad dollars spent, then tune for precision

$$F1 = 2 \frac{\text{Recall} * \text{Precision}}{\text{Recall} + \text{Precision}}$$

Overfitting and Underfitting



$$P(\hat{y} = 1|x) = \sigma(w_1 * x_1 + w_2 * x_2 + b)$$

**A more
Complex Model**

$$P(\hat{y} = 1|x) = \sigma(w_1x_1 + w_2x_2 + w_3x_1x_2 + w_4x_1^2 + w_5x_2^2 + b)$$

A simpler model

$$P(\hat{y} = 1|x) = \sigma(w_1x_1 + b)$$

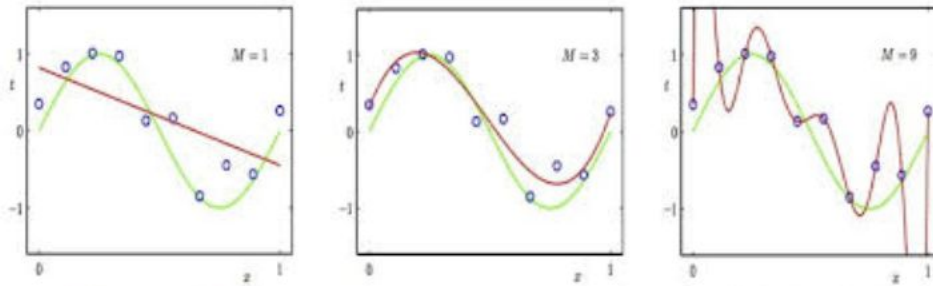
**A highly simplistic
model**

$$P(\hat{y} = 1|x) = \sigma(b)$$

Model Complexity

Under- and Over-fitting examples

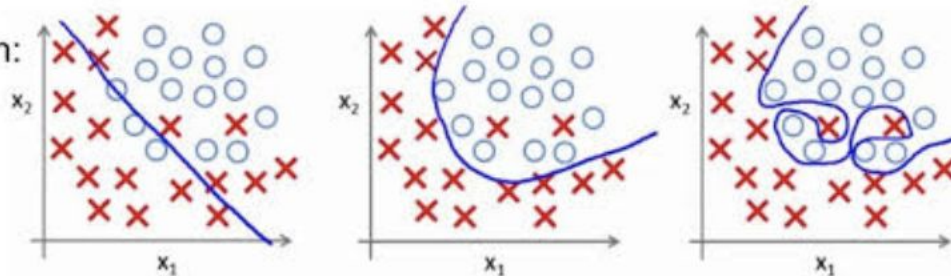
Regression:



predictor too inflexible:
cannot capture pattern

predictor too flexible:
fits noise in the data

Classification:



Copyright © 2014 Victor Lavytchenko

Underfit Models:

- Model not complex enough to capture the underlying distribution of the data.

Overfit Models:

- Model too complex and tries to capture every bit of information in the dataset.
- Such models do not generalize well to unseen data.

Model Complexity in the case of Logistic Regression could be due to:

- Large number of parameters.

Fixing Underfitting

- Add more features to your Logistic Regression Model
- Try using a more complex model, such as :
 - Decision Trees
 - Neural Nets

Fixing Overfitting

- Collect more data
- Then reduce model complexity
 - Try **regularization**
 - Then try a simpler model

QnA & Code Walkthrough

Scikit Learn Cheat Sheets

- https://scikit-learn.org/stable/tutorial/machine_learning_map/
- https://s3.amazonaws.com/assets.datacamp.com/blog_assets/Scikit_Learn_Cheat_Sheet_Python.pdf
- <https://bit.ly/2Kwg36X>
- <https://towardsdatascience.com/resources-to-start-your-journey-in-data-science-bf960a8d928c>

Thank you

Supervised Machine Learning

